


Multi-layer Perceptron


Artificial Intelligence

School of Computer Science
The University of Adelaide

Recall Single-layer Perceptron



XOR




Single layer perceptrons are linear classifiers.

Perceptron algorithm will not converge for linearly non-separable problems.

Multi-layer Perceptron

A Non-linear Classifier




MLPs are more expressive than Perceptrons since they can learn highly non-linear class boundaries.

Multi-layer Perceptron

A Non-linear Classifier



The most common case involves a single hidden layer:



Multi-layer Perceptron

A Non-linear Classifier


Each *hidden unit* can be considered as single output perceptron network:



Multi-layer Perceptron

A Non-linear Classifier


Which is capable of separating the training examples linearly



Multi-layer Perceptron

A Non-linear Classifier


The output unit of the multi-layer network can then be considered a soft-thresholded linear combination of the hidden units (which are equivalent to single output unit perceptrons):



Multi-layer Perceptron


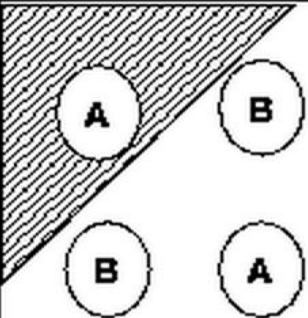
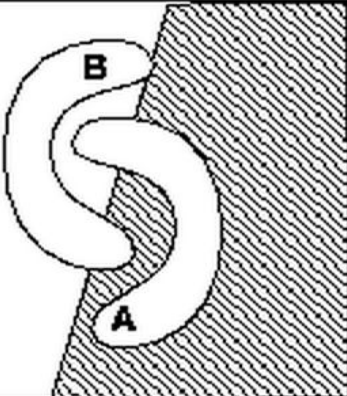
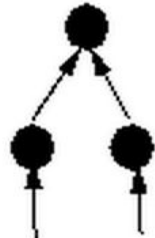
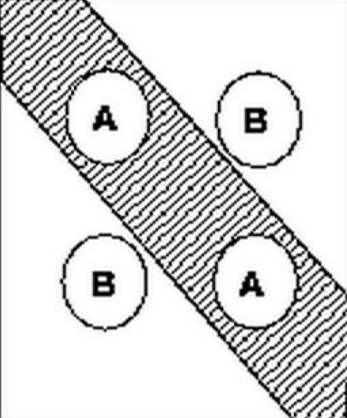
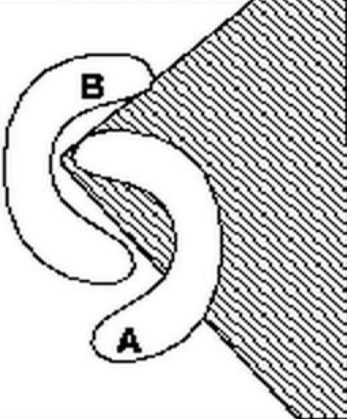
A Nonlinear Classifier

Remember the **AND** and **OR** function implementation using a single artificial neuron?




Multi-layer Perceptron

A Nonlinear Classifier


Structure	XOR	Meshed regions
single layer 		
two layer 		

Multi-layer Perceptron

A Nonlinear Classifier



(a)



(b)

Figure 20.23 (a) The result of combining two opposite-facing soft threshold functions to produce a ridge. (b) The result of combining two ridges to produce a bump.


Multi-layer Perceptron

The Good and the Bad

- With a single, sufficiently large hidden layer, it is possible to represent *any continuous* function of the inputs with *arbitrary* accuracy.
- Unfortunately, for any *particular* network, it is harder to characterize exactly which functions can be represented and which ones cannot.
- As a consequence, given a particular learning problem, it is unknown how to choose the *right number of hidden units* in advance.
- One usually resorts to cross validation, but this can be computationally expensive for large networks.

Backpropagation

Understanding Perceptron Training



- The function that a perceptron network corresponds to can be represented as $h_{\mathbf{W}}(\mathbf{x})$, where

$$h_{\mathbf{W}}(\mathbf{x}) = g(inputs) = g\left(\sum_{j=0}^D W_j x_j\right)$$

Understanding Perceptron Training

- Perceptron learning (generally, neural network learning) occurs by *adjusting the weights to minimize some measure of error*.
- Let (\mathbf{x}, y) be a *single* training sample with its *true* output y . The squared error is given by

$$\begin{aligned} E &= \frac{1}{2} Err^2 \\ &= \frac{1}{2} (y - h_{\mathbf{W}}(\mathbf{x}))^2 \\ &= \frac{1}{2} (y - g(\sum_{j=0}^D W_j x_j))^2 \end{aligned}$$

Note scaling the error with $\frac{1}{2}$ does not change its minimizer.

Understanding Perceptron Training

- Calculating the partial derivative of the error against a particular weight, we have

$$\begin{aligned}\frac{\partial E}{\partial W_j} &= Err \times \frac{\partial Err}{\partial W_j} \\ &= Err \times \frac{\partial}{\partial W_j} \left(y - g\left(\sum_{j=0}^D W_j x_j\right) \right) \\ &= -Err \times g'\left(\sum_{j=0}^D W_j x_j\right) \times x_j\end{aligned}$$


where g' is the derivative of the activation function g .

Understanding Perceptron Training

- Under the **gradient descent** algorithm, if we want to *reduce* E , we update the weight as follows:

$$W_j \longleftarrow W_j + \alpha \times Err \times g'(\sum_{j=0}^D W_j x_j) \times x_j$$

where α is the **learning rate**.



Multi-layer Perceptron Training Backpropagation

(derivation on the board)

Backpropagation

Multi-output Multi-layer Perceptron

- We need to consider multiple output units for multi-layer networks. Let (\mathbf{x}, \mathbf{y}) be a single sample with its desired output labels $\mathbf{y} = \{y_1, \dots, y_i, \dots, y_M\}$.
- The error at the output units is just $\mathbf{y} - h_{\mathbf{W}}(\mathbf{x})$, and we can use this to adjust the weights between the hidden layer and the output layer.
- The above steps produces a term equivalent to the error at the hidden layer, i.e. the error at the output layer is **back-propagated** to the hidden later.
- This is subsequently used to update the weights between the input units and the hidden layer.

Backpropagation

Step 1: Update the weights between the hidden and output layers.

- Let Err_i be the i -th component of the error vector $\mathbf{y} - h_{\mathbf{W}}(\mathbf{x})$.
- Define $\Delta_i = Err_i \times g'(in_i)$.
- The weight-update rule becomes

$$W_{j,i} \longleftarrow W_{j,i} + \alpha \times a_j \times \Delta_i$$

This is similar to weight-updates for Perceptrons!

Backpropagation

Step 2: Back-propagate the error to the hidden layer.

- The idea is that the hidden node j is “responsible” for some fraction of the error Δ_i in each of the output nodes to which it connects.
- Thus the Δ_i values are divided according to the strength (weight) of the connection between the hidden node and the output node:

$$\Delta_j = g'(in_j) \sum_i W_{j,i} \Delta_i$$

Backpropagation

Step 3: Update the weights between the input units and the hidden layer.

- Again, this is similar to weight-updates in Perceptrons:

$$W_{k,j} \longrightarrow W_{k,j} + \alpha \times a_k \times \Delta_j$$

Backpropagation

For the general case of *multiple hidden layers*:

- ① Compute the Δ values for the output units, using the observed error.
- ② Starting with the output layer, repeat the following for each layer in the network, until the earliest hidden layer is reached:
 - Propagate the Δ values back to the previous layer.
 - Update the weights between the two layers.
- ③ Repeat Steps 1 to 2 for all training samples.

Backpropagation


```
function BACK-PROP-LEARNING(examples, network) returns a neural network
  inputs: examples, a set of examples, each with input vector x and output vector y
           network, a multilayer network with  $L$  layers, weights  $W_{j,i}$ , activation function  $g$ 

  repeat
    for each  $e$  in examples do
      for each node  $j$  in the input layer do  $a_j \leftarrow x_j[e]$ 
      for  $\ell = 2$  to  $M$  do
         $in_i \leftarrow \sum_j W_{j,i} a_j$ 
         $a_i \leftarrow g(in_i)$ 
      for each node  $i$  in the output layer do
         $\Delta_i \leftarrow g'(in_i) \times (y_i[e] - a_i)$ 
      for  $\ell = M - 1$  to  $1$  do
        for each node  $j$  in layer  $\ell$  do
           $\Delta_j \leftarrow g'(in_j) \sum_i W_{j,i} \Delta_i$ 
          for each node  $i$  in layer  $\ell + 1$  do
             $W_{j,i} \leftarrow W_{j,i} + \alpha \times a_j \times \Delta_i$ 
  until some stopping criterion is satisfied
  return NEURAL-NET-HYPOTHESIS(network)
```

Figure 20.25 The back-propagation algorithm for learning in multilayer networks.

Backpropagation with SGD

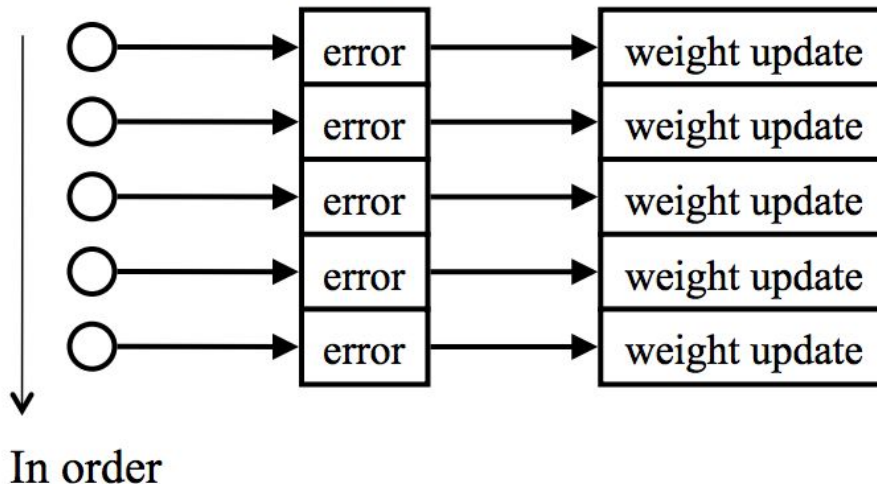
Batch: training step



Gradient Decent

Batch: training over all given examples once.

Online: training step


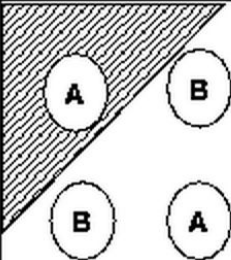
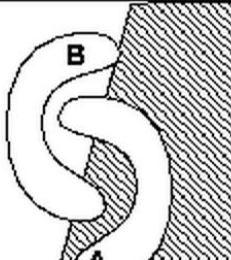

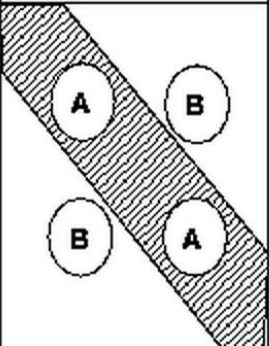
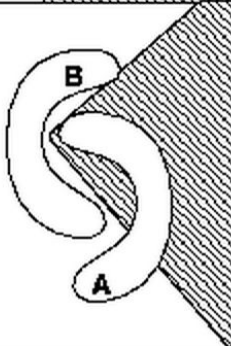
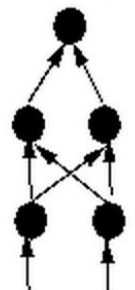
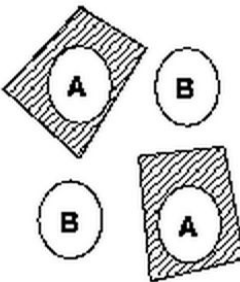
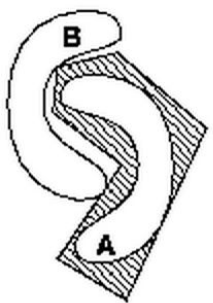


Stochastic Gradient Decent



- Randomly choose m samples
- Compute the gradients
- Do backpropagation
- Repeat


Multi-layer Perceptron

A Nonlinear Classifier

Structure	Regions	XOR	Meshed regions
single layer 	Half plane bounded by hyper- plane		
two layer 	Convex open or closed regions		
three layer 	Arbitrary (limited by # of nodes)		

Deep Multi-layer Perceptron

- We can learn anything !!!
- More than one hidden layer  **deep.**
- Higher level representations  Better at high-level tasks.
- Visual Classification / Speech Recognition / Scene understanding / Visual question answering



Not so fast... 😞

- Backpropagation [late 80s, early 90s]
 - Goal was to train nets with large number of layers, so that features could be learned directly from input data, but it didn't quite work
 - Notable exception: convolutional neural net by Y. LeCun (large # layers, but small # parameters)
- Issues with MLP training via backpropagation
 - Very slow convergence, particularly in large nets and large databases
 - Slow computers of the 80s and 90s
 - Local minima (how to initialize SGD)
 - Net structure (cross validation)
 - Overfitting
- In the 90s and early 2000
 - development of several classifiers (e.g., Boosting, SVMs)
 - hand-designed hierarchical representations (e.g., bag of features)