

# Deep Learning

Artificial Intelligence

School of Computer Science  
The University of Adelaide

# Visual Learning

- Learning the mapping function.

$$f(\mathbf{X}) \rightarrow \mathbf{Y}$$

- For example, Image to category, video to action



This image by Nalla is  
licensed under CC-BY 2.0

(assume given set of discrete labels)  
{dog, cat, truck, plane, ...}

→ cat

# Visual Learning

- Learning the mapping function.

$$f(\mathbf{X}) \rightarrow \mathbf{Y}$$

- **Why is it difficult?**



This image by Nalla is  
licensed under CC-BY 2.0

(assume given set of discrete labels)  
{dog, cat, truck, plane, ...}

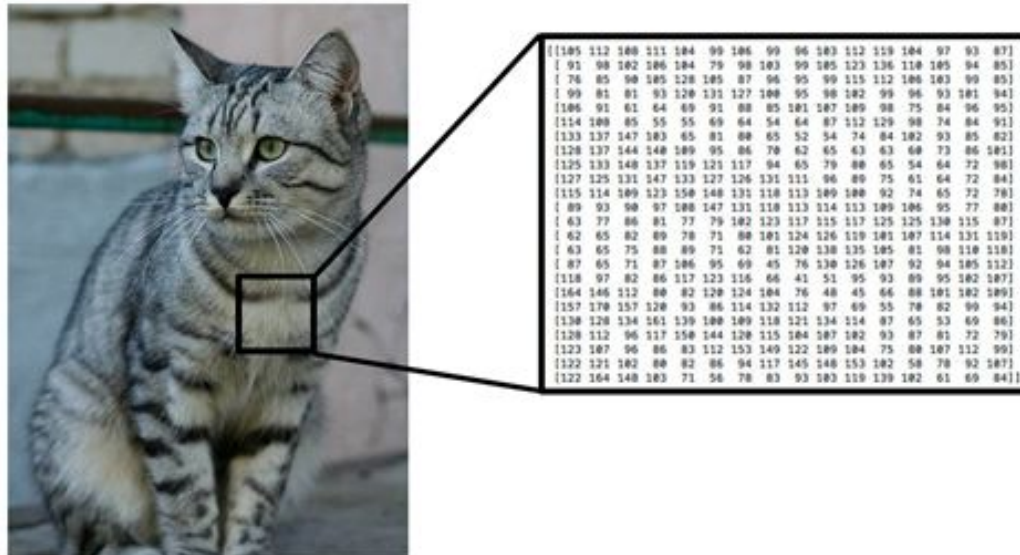
→ cat

# Visual Learning

- Learning the mapping function.

$$f(\mathbf{X}) \rightarrow \mathbf{Y}$$

- Why is it difficult?

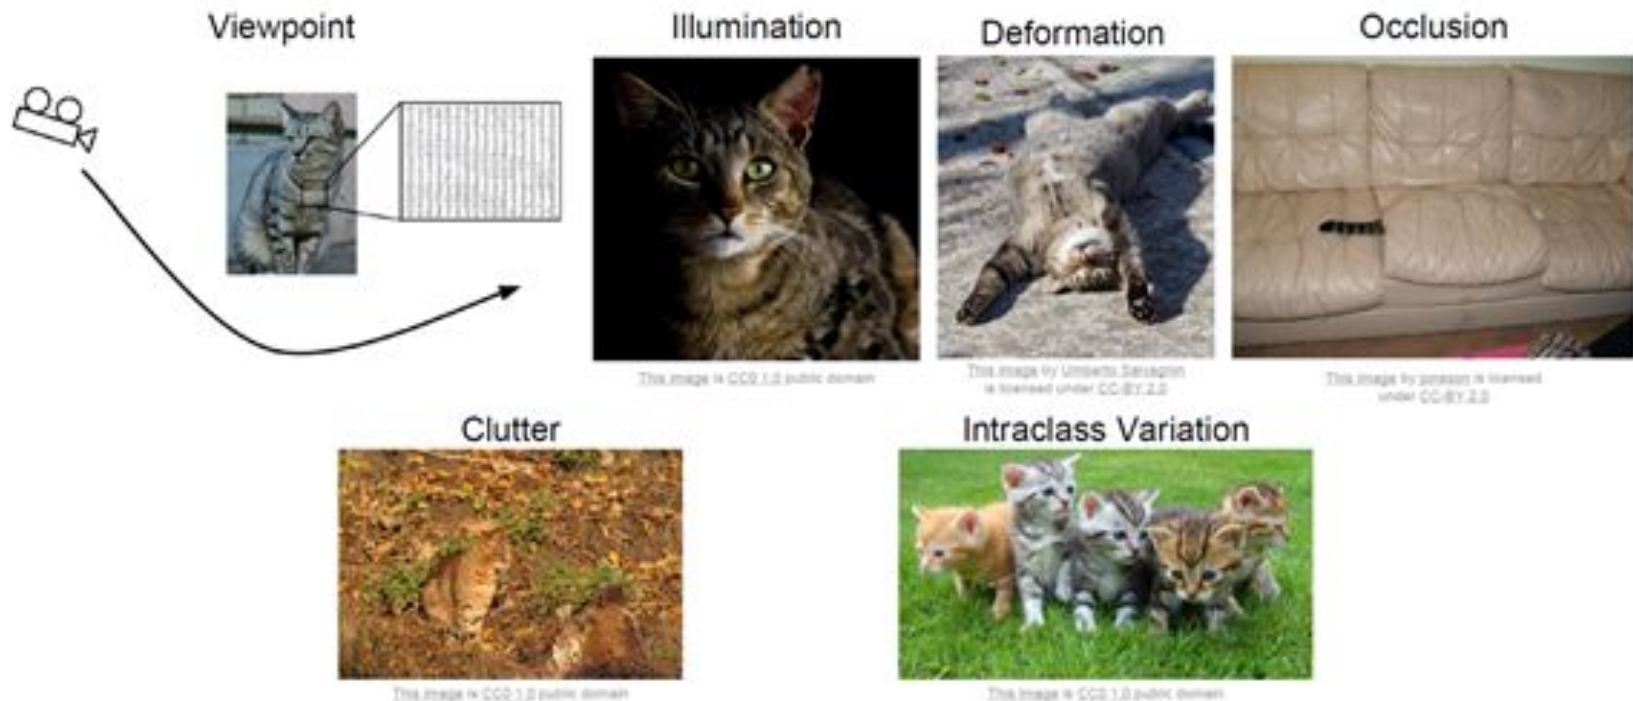


# Visual Learning

- Learning the mapping function.

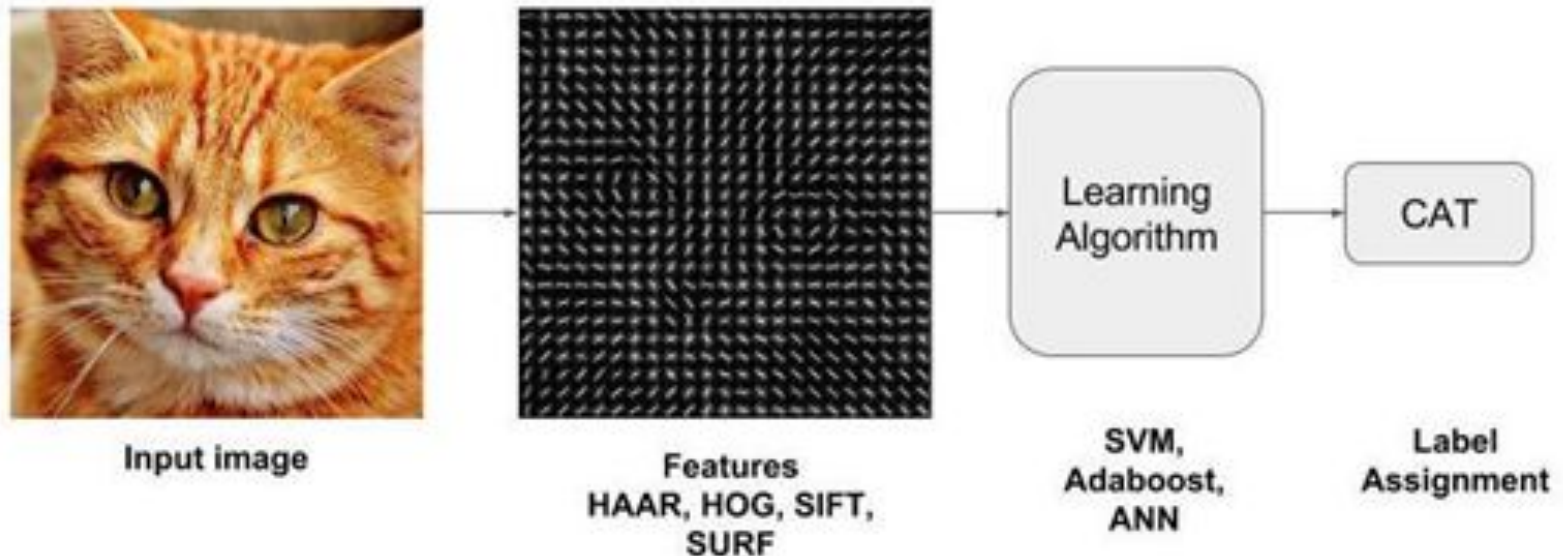
$$f(\mathbf{X}) \rightarrow \mathbf{Y}$$

- Why is it difficult?



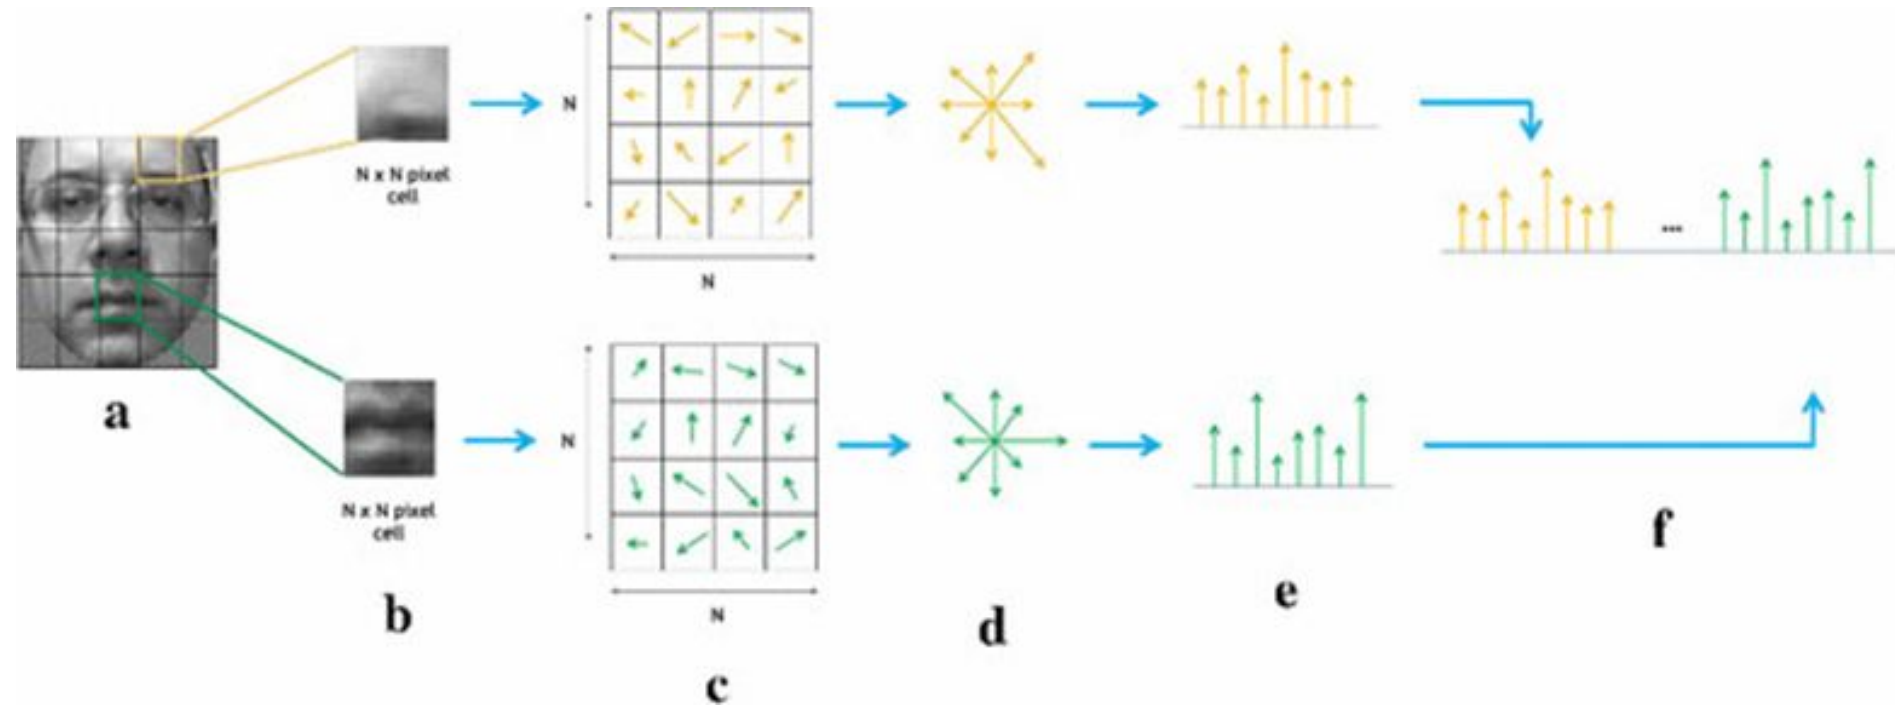
# Traditional Learning

- Decompose the problem into two parts:



- Hand crafting features and learning a classifier.

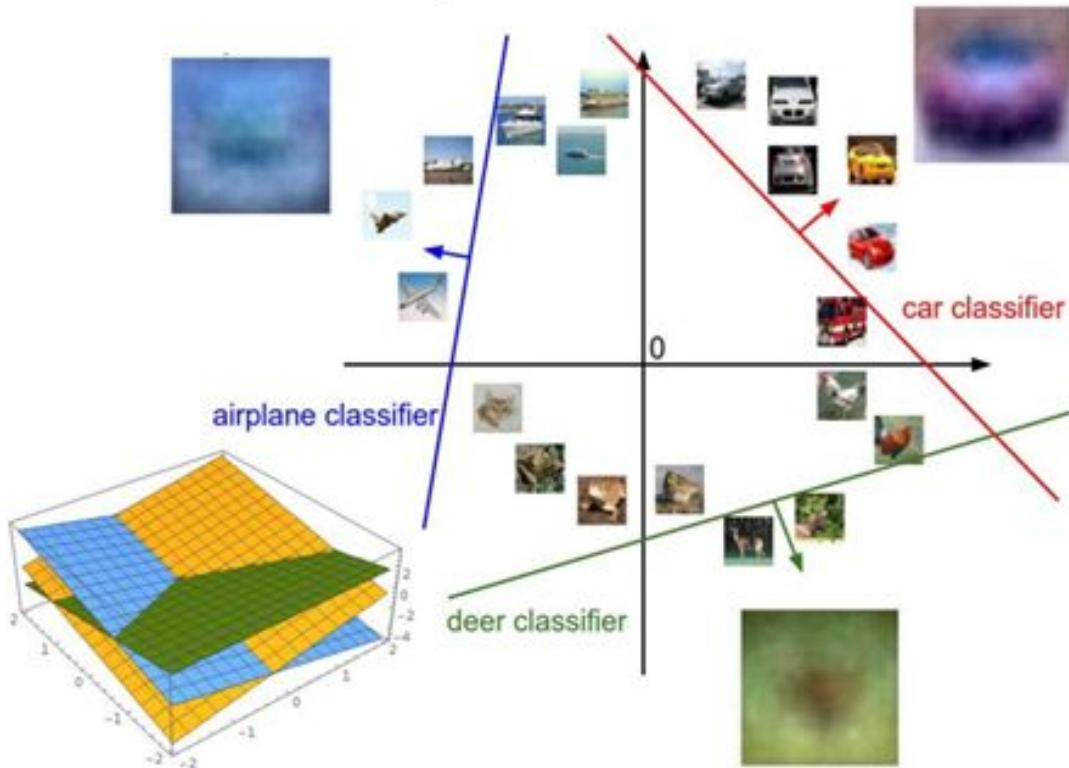
# Example of Hand-Crafted Feature (HOG)



- Engineering features with desired invariances is difficult.
- You can easily lose important information.



# Example of Linear Classifier



$$f(x, W) = Wx + b$$



Array of **32x32x3** numbers  
(3072 numbers total)

- Classification = learning parameters  $W$  and  $b$ .
- What shape is the classification boundary?
- How to optimize?



# Deep Network

- Decompose the problem into multiple parts:


$$\mathbf{Z}_1 = f_1(\mathbf{X})$$

$$\mathbf{Z}_2 = f_2(\mathbf{Z}_1)$$

...

$$\mathbf{Z}_t = f_t(\mathbf{Z}_{t-1})$$

$$\mathbf{Y} = f_y(\mathbf{Z}_t)$$



- Both features and classifier are learned

# Deep Network

- Decompose the problem into multiple parts:


$$\mathbf{Z}_1 = f_1(\mathbf{X})$$

$$\mathbf{Z}_2 = f_2(\mathbf{Z}_1)$$

...

$$\mathbf{Z}_t = f_t(\mathbf{Z}_{t-1})$$

$$\mathbf{Y} = f_y(\mathbf{Z}_t)$$



- Both features and classifier are learned
- Researchers have hypothesized that the number of layers in an MLP correlates well with high-level information.

# MLP Based Deep Learning


- Unmanageable number of parameters
  - Example, let's take MNIST problem with  $28 \times 28 = 784$  input images
  - 1 hidden layer with 1000 nodes and an output layer with 10 nodes: # weights =  $784 \times 1000 + 1000 \times 10 = 794000$  weights
  - 2 hidden layers with 1000 nodes and an output layer with 10 nodes: # weights =  $784 \times 1000 + 1000 \times 1000 + 1000 \times 10 = 1794000$  weights
  - Prone to overfitting
- Gradient descent didn't work beyond a couple of hidden layers
  - Magnitude kept reducing as the gradient flowed back to the input layer (vanishing gradient problem [Hochreiter91])
  - Convergence issues
- Machines (at the time – 80s and 90s) couldn't cope with datasets bigger than a few thousand samples and models with more than a few thousand weights

# Learning Neural Networks

- Optimizing a loss function to learn parameters

$$L(W) = \frac{1}{N} \sum_{i=1}^N L_i(f(x_i, W), y_i)$$

Fitting to data



**Too many  
Parameters  
=  
Overfitting!!!**

# Regularization


- Optimizing a loss function to learn parameters

$$L(W) = \underbrace{\frac{1}{N} \sum_{i=1}^N L_i(f(x_i, W), y_i)}_{\text{Fitting to data}} + \underbrace{\lambda R(W)}_{\text{Choose the simplest model}}$$

## Overfitting

Remember Occam's  
Razor !!!

Usually L2 or L1 sum of  
the network's weights are  
minimized to select  
simplest hypothesis.






# Learning of Neural Network

- Optimizing a loss function to learn parameters

$$L(W) = \underbrace{\frac{1}{N} \sum_{i=1}^N L_i(f(x_i, W), y_i)}_{\text{Fitting to data}} + \underbrace{\lambda R(W)}_{\text{Choose the simplest model}}$$

Fitting to data

Choose the simplest model




## Overfitting

Remember Occam's  
Razor !!!

Regularizing decision  
boundaries in this manner  
help avoiding overfitting.


# Non-Linear Activations



$$\mathbf{h} = s(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$$
$$\hat{\mathbf{y}} = \mathbf{W}_2 \mathbf{h} + \mathbf{b}_2$$




# Non Linearities are Important



# Vanishing Gradient Problem


- Non-linearity based on sigmoid.



$$g_{sig}(in) = \frac{1}{1 + e^{-in}}$$

# Vanishing Gradient Problem


- Non-linearity based on sigmoid.



Small gradient magnitude, particularly at the tails

# Vanishing Gradient Problem


- After several multiplications over the layers, the grad mag will become insignificant.



Small gradient magnitude, particularly at the tails

# Vanishing Gradient Problem


- Note that we could try to fix the problem by using non-linear activations that have gradients of larger magnitude
  - But if the magnitude gets larger than 1, then we have the problem of exploding gradients



# Rectified Linear Units (RELU)

Nair & Hinton (2010)


- Maximum gradient magnitude is 1
- Large region where this magnitude is 1
- Still non-linear
- Gradient shape?



# Some Prevalent Activation Functions


## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




## tanh

$$\tanh(x)$$




## ReLU

$$\max(0, x)$$



## Leaky ReLU

$$\max(0.1x, x)$$




## Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

## ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$





# Biggest Influence in Visual Perception Change in NN Architecture Deep Convolutional Neural Networks

